



PRODUCT OVERVIEW • JUNE 2026

A SECURE COMPUTE PLATFORM FOR AI.

Product overview for CyberKnight technical review.

MIKE GAULT • CEO • mike@unity-labs.com
ALAN RADI • COMMERCIAL • alan@unity-labs.com

THE FOURTH REBUILD OF COMPUTING.

Mainframe → Client-server → Cloud → Mobile → **Agentic**

Identity. Execution. Governance. Payments.

All need to be rebuilt for AI Agents without a human in the loop.

MARKET OPPORTUNITY

MARKET OPPORTUNITY: A TRILLION-DOLLAR REBUILD.

NVIDIA

"In ten years, we will have 75,000 employees working with 7.5 million agents."

– Jensen Huang, Nvidia GTC, March 2026

100:1

agents-to-humans ratio at scale

GARTNER

"By 2028, the average Fortune 500 will run 150,000+ AI agents, up from fewer than 15 today. Only 13% have adequate governance."

– Max Goss, Gartner, April 2026

10,000×

agent growth in 3 years · 87% ungoverned

STRIPE

"We will likely need blockchains that support more than one million – or even one billion – transactions per second."

– Patrick & John Collison, Stripe, 2026

1B TPS

the infrastructure thesis, validated




If we want an internet where AI agents can operate without destroying trust, **blockchains are not optional infrastructure**. They are the missing layer that makes an AI-native internet work.

– **A16Z CRYPTO, FEBRUARY 2026**

THE PLATFORM

THE SECURE AI COMPUTE PLATFORM

Every requirement autonomy demands, delivered as one protocol stack.



EVERY AUTONOMY REQUIREMENT, DELIVERED

✓ IDENTITY · ✓ EXECUTION · ✓ GOVERNANCE · ✓ PAYMENTS

AGENT FRAMEWORKS

THE APPLICATION LAYER

Bring your own agent – any framework, any model.

LANGCHAIN · OPENCLAW · AUTOGEN · CUSTOM

UNICITY AOS

RUNTIME

Executes the agents. Capsules provide multi-tenancy and isolation. Every prompt, tool call, and decision passes through it – proactive safety policy, goal focus and budget enforcement, with tokenized execution trace for analysis and compliance, with hard policy rules on tool interactions.

KERNEL · POLICY · INTERCEPT


UNICITY PROOF SYSTEM

PROOF SYSTEM · FORMAL VERIFICATION

State tokenization. Verifiable execution. Cryptographic identity. Value transfer.

AOS KERNEL ARCHITECTURE.

AOS is the secure enforcement kernel that tools like Claude Code and OpenClaw run **on top of**, not alongside. A user-space microkernel with syscalls, airlocks, and sandboxed isolation – the kernel stays minimal while everything else lives in user-space.



User-Space Microkernel

AOS models itself as an OS – with syscalls, airlocks, and user-space isolation. The kernel (aos9d) and system SDK (aos9-sys) are the only core components.

WASM Capsule Sandbox

All plugins run as WASM capsules inside a sandboxed runtime. OpenClaw plugins are compiled into capsules via aos9-openclaw, with host calls thunked into aos9::sys.

Enforcement Layer


Security interception, budget enforcement, and audit logging happen at the runtime level – below the agent, not beside it. Tools cannot bypass the kernel.

Everything Is User-Space

CLIs, OpenClaw, frontends, and even LLM providers are external user-space components communicating over IPC. The kernel stays minimal and ignorant of legacy ABIs.

EVERYONE ELSE WATCHES THE AGENT. WE RUN IT AND PROVE ITS EXECUTION.

Verifiable execution and cryptographic guarantees of security – not behavioral inference, not log signals. **Proof.**



A O S A R C H I T E C T U R E

THE SEMANTIC FIREWALL

The **Semantic Intercept Fabric (SIF)** goes beyond action-based policies to content-based controls with domain-specific ontologies – a Rust implementation with optimized search and rule algorithms.

→ **WASM SANDBOXING**

Agents execute within a WebAssembly isolated container. No direct access to the host OS.

→ **IPC MESSAGE BUS**

All external communication is routed through a monitored inter-process communication bus, enabling continuous, tamper-evident auditing.

→ **SUB-20MS ENFORCEMENT**

Intent is classified inline and enforced before the action executes – sub-twenty-millisecond overhead, in-path, not after the fact.

1 · RULE ENGINE

DETERMINISTIC

High-performance pattern matching and candidate identification. Enables dynamic rule updates to catch zero-day injection vectors without restarting the cluster.

2 · ML DETECTION

INFERENCE

Low-latency inference and anomaly detection. Plugin-expandable, ONNX-compatible model loading allows CISOs to deploy custom detection models.


3 · DLP SCANNER

EXFILTRATION

Continuous scanning for secrets, IP, and PII exfiltration using advanced pattern matching and Named Entity Recognition (NER).

DEPLOYMENT

ENTERPRISE DEPLOYMENT.



PRODUCT

ENTERPRISE CASE STUDIES.

Unicity's secure compute platform delivers what enterprises need today – from DLP to cost control and developer freedom.

- 01

INTELLIGENT ROUTING

Route every agent request to the right model, tool, or capsule.
- 02

BRING YOUR OWN CODING AGENT

Organizational freedom for any framework with security and strict cost control.
- 03

AGENT SPRAWL COST CONTROL


Budget enforcement and token accounting per agent and per workflow.
- 04

DATA LOSS PREVENTION

Semantic intercept fabric blocks exfiltration before it leaves the perimeter.
- 05

X402 PAYMENTS

Private, compliant A2A settlement over Unicity Proof System.



Live: Unicity operator console – one pane of glass across every agent workload.

USE CASE · VISIBILITY

AGENT SPRAWL & SHADOW AI

You cannot secure what you cannot see. Unicity acts as the mandatory registration and monitoring layer for all machine identity in the enterprise.

→ CRYPTOGRAPHIC IDENTITY

Every agent is assigned a revocable identity bound to a human principal. Eliminates static API keys.

→ GLOBAL INVENTORY

Map every active agent and its authorized scopes across the organization.

→ DRIFT CONTAINMENT

Instantly quarantine rogue agents that deviate from baseline execution patterns.

AGENT SPHERE

acme-industrial

OverviewSecurityGovernanceNetworkingPaymentsAuditDevelopers

prod

Search...

⌵

MC

Overview / Governance

G

GOVERNANCE

Identity, policy, approvals, and audit. Define what agents can do, gate the actions that need a human, and keep a sealed record of every decision.

Last 24 hoursExport audit bundle+ New policy

ACTIVE POLICIES

70

+4 this week

APPROVALS WAITING

5

12 min oldest

DECISIONS (24H)

185k

98.5% allow

AUDIT HASHES SEALED

185,144

SHA-256 chained · 7y

Decisions by category · 24h

185,144 total

Security		31,488
Identity		42,008
Spend		12,884
Routing		68,442
PII / data		28,114
Approval		2,208

Allow 182,408Review 1,489Deny 1,247

Approvals waiting

5 items

\$

finance-ops · transfer \$5,200 → Bank of FR

policy: requires-human-approval · transfers > \$1,000

12 min

Approve

Reject

↶

cs-bot · refund \$89 → customer #42811

policy: cs-refund-any-amount · > \$50

4 min

Approve

Reject

✓

procurement · commit RFQ to FreightLink AI · \$28,400

policy: vendor-commit-cap · > \$25k

1 min

Approve

Reject

⌚

sales-assist · open lane → SF Agent Cloud

policy: new-counterparty · first-time vendor

38 min

Approve

Reject

⌚

eng-copilot · model-version upgrade

policy: model-allowlist · new model release

2 hr

Approve

Reject

USE CASE · DATA LOSS PREVENTION

DATA LOSS PREVENTION (DLP).

Unicity sits directly in the transport flow, inspecting LLM payloads before they ever leave the corporate perimeter to prevent exfiltration.

→ **INLINE REDACTION**


Detect and mask PII, credentials, and corporate secrets in the prompt before inference.

→ **INJECTION DEFENSE**

Blocks agents manipulated into leaking data via indirect prompt injections.

→ **COMPLIANCE LOGS**

Cryptographic, tamper-evident logs of all data touched by agents for compliance audits.



USE CASE · COST CONTROL

RUNAWAY COMPUTE & COST CONTROL

A poorly engineered agent stuck in a hallucination cycle can burn thousands in token costs. We provide the financial circuit breakers.

→ HARD RATE LIMITING


Enforce strict execution thresholds at the kernel level before cost is incurred.

→ TOKEN QUOTAS

Assign budget limits. The AOS Kernel blocks execution instantly when reached.

→ ROI ATTRIBUTION


Track exactly how much compute each autonomous workflow consumes.



— CAPABILITY · ORCHESTRATING THE AGENT MESH

INTELLIGENT ROUTING.

Every agent request is scored before it runs and directed to the right model – by task complexity, cost, latency, and security clearance. The routing decision is itself an inspected, recorded event, so traffic is matched to capability without ever escaping the security boundary.



— **CAPABILITY-MATCHED**

Heavy work reaches strong models; trivial work never pays for them.

— **CLEARANCE-AWARE**

Sensitive context is only ever routed to a target authorised to receive it.

— **RECORDED**

The routing choice is an inspected event – every decision is explainable after the fact.

THE TEAM

WHO ARE WE?

Unicity Labs. A team of PhDs in machine learning and cryptography with deep experience building agentic systems – applying fifteen years of nation-state-grade cryptographic engineering to the rail your agents transact on.



Mike Gault

CEO

- PhD Electrical Engineering
- Built & exited Guardtime (ADX:IHC)
- Ex-MD, Barclays Capital



Tony Kenyon

CPTO

- PhD Machine Learning
- 25 years shipping enterprise AI & infra
- Principal Architect: BT, Nokia, A10




Alan Radi

COMMERCIAL

- 12 years implementing CX AI for Global B2C brands
- Apple · Google · Verizon · Pepsi · DHL

WHERE WE OPERATE



CRYPTOGRAPHIC INFRASTRUCTURE TRUSTED BY



SEE IT IN ACTION.


Book a Live Demo

See an agent execute under policy, with every decision verified and audit-logged – end-to-end.

 info@unicity.ai

 unicity.ai

 +971 50 884 3756



Unicity Labs

502, 5th Floor, Al Montazah Tower B
Khalidiyah Area, Abu Dhabi, UAE

THANK YOU