



PLATFORM OVERVIEW • 2026

THE SECURE AGENTIC COMPUTE PLATFORM **FOR** **HEALTHCARE.**

MIKE GAULT • CEO • mike@unicity-labs.com

WHY NOW

THE MAGIC AGENT HAS ARRIVED.

BOOKING

MESSAGING

TRIAGE

FOLLOW-UP

AVAILABLE 24/7

A swarm of vendors is promising an agent that just handles the work – booking, messaging, triage, follow-up, around the clock. It's a genuinely compelling pitch.

BUT THIS WORK TOUCHES DATA, RECORDS, AND PATIENTS.

THE TRAP

THE MOST REGULATED BEHAVIORS ON EARTH.

WHAT THE AGENT DOES

Stores a photo

Quotes a price

Books a slot

Messages a patient

Drifts toward a diagnosis

WHAT IT LEGALLY IS

→ **PHI PROCESSING & CONSENT**

→ **COMMERCIAL & ADVERTISING RULES**


→ **PATIENT DATA HANDLING**

→ **COMMUNICATION & CONSENT**

→ **PRACTICING MEDICINE**

The capability you're buying and the liability you're taking on are **the same five actions.**

IT'S 2AM. NO ONE IS ON THE LINE.



- 1 A patient messages the agent at 2AM. No human is on the line.
- 2 The conversation runs long. The only boundary is a prompt.
- 3 Rapport builds, symptom detail deepens, the model reaches hallucinated confidence.
- 4 The agent issues a prescription.

ONE MOMENT • THREE BREACHES

Prescribing without a license

PHI stored without consent

An illegal cross-border data transfer

One conversation. Three regulatory regimes breached. Drawn by prompt engineering – which a clever patient, an edge case, or hallucinated confidence will eventually break.

THE COST OF FAILURE

FOUR REGIMES. EACH WITH TEETH.

AI RULES

EU AI ACT

- High-risk classification
- Core obligations from Aug 2026
- Deployer duties, not just providers

PRIVACY

GDPR

- Explicit consent required
- Data minimization
- Auditable data lifecycle

SOVEREIGNTY

DATA RESIDENCY

- Where data physically lives
- Cross-border transfer limits
- Provable in-region processing

MEDICAL

LICENSING

- Practicing medicine unlicensed
- Diagnosis & prescription
- Personal & criminal liability

These regimes **compound** – one failure triggers several at once. Penalties stack.

~26% of hospital representatives feel ready for EU AI Act obligations.

THE OPPORTUNITY

PROMPTS AREN'T A BOUNDARY. KERNELS ARE.

GUARDRAILS IN THE PROMPT

BYPASSABLE

A clever patient, a strange edge case, or hallucinated confidence will eventually talk its way around a sentence of instructions.

ENFORCEMENT IN THE KERNEL

UNBYPASSABLE


Every regulated action is intercepted in-path by an enforcement kernel the agent cannot talk around. It cannot break the kernel.

The villain isn't the Magic Agent – it's that its guardrails live in prompt engineering. Unicity is the layer the agent **runs on**: bring any agent, any framework, any vendor – **we make it safe to deploy**.

THE PLATFORM

UNICITY: THE SECURE AI COMPUTE PLATFORM.

Bring any agent. Every regulated clinical action passes through one enforcement stack.



— AGENT FRAMEWORKS

THE APPLICATION LAYER

Bring your own agent – any framework, any model.

LANGCHAIN · OPENCLAW · AUTOGEN · CUSTOM

— UNICITY AOS

MULTI-TENANT SECURE OS

Executes the agents. Capsules provide multi-tenancy and isolation. Every prompt, tool call, and clinical decision passes through it – regulated actions are intercepted in-path, with hard policy rules and a tokenized execution trace for audit and compliance.

KERNEL · POLICY · INTERCEPT

— UNICITY BLOCKCHAIN

PROOF SYSTEM · FORMAL VERIFICATION

State tokenization. Verifiable execution. Cryptographic identity. Value transfer.


EVERY AUTONOMY REQUIREMENT, DELIVERED

✓ IDENTITY · ✓ EXECUTION · ✓ GOVERNANCE · ✓ COMPLIANCE

DEPLOYMENT

RUNS INSIDE YOUR INFRASTRUCTURE.

The agent runs inside your network, behind a single egress gate. **Patient data and the clinical moment never leave your network** – the GDPR, sovereignty, and cross-border-transfer answer, all at once.



WHY WE'RE DIFFERENT

BUILT TO RUN AGENTS AT SCALE — SECURELY, EFFICIENTLY, PROVABLY.

Unicity is the first secure, efficient, and provable agent compute platform — the layer a hospital's compliance officer signs off on before an agent touches a patient.

01 · SECURELY

Prompt-proof boundaries

The agent runs inside your kernel, on your infrastructure. Every regulated action is intercepted in-path — security can't be bypassed, and data never leaves your network.

AOS enforcement kernel · in-path, unbypassable

02 · EFFICIENTLY

A price point that works at scale

Intelligent routing sends each task to the cheapest model that can handle it, and multi-tenancy packs many agents onto the same hardware. Together they bring the cost per agent low enough to run one on every line — thin hospital margins included.

Intelligent routing + multi-tenancy · cost per agent that scales

03 · PROVABLY

Provable, not just observed

Every regulated action lands in a tamper-evident record. You don't infer what an agent did from scattered logs — you have one definitive, audit-ready answer. Exactly what an EU AI Act deployer is legally required to hold.

Tamper-evident audit · AI Act deployer-ready

THE PAYOFF · ONCE IT'S CONTAINED

NOW YOU CAN **ACTUALLY DEPLOY IT.**

Containment is what **unlocks** the upside. Once the agent is safe to run, the same platform delivers the documentation-burden and clinician-retention wins.

+30.7%

Documentation-related wellbeing
after ambient documentation

EMORY HEALTHCARE

-21.2%

Burnout prevalence after 84
days

MASS GENERAL BRIGHAM

-27%

Time spent on notes per
appointment

INTERMOUNTAIN

-41%

Physician clerical time in live
autonomous-documentation
deployments

INDUSTRY DEPLOYMENTS

Replacing a single physician costs **~\$500K-\$1M**. Cutting documentation burden is a retention strategy with a hard financial case.


The clinician stays the final arbiter – the platform **enforces** that, and every AI action is source-linked and verifiable.

ONE PANE OF GLASS. EVERY REGULATED ACTION.

01 SECURITY

INTERCEPT

Every regulated action stopped in-path before it happens.



02 GOVERNANCE

POLICIES · APPROVALS · SEALED AUDIT


Who can do what, human sign-off, and a tamper-evident record.



03 NETWORKING

CROSS-BORDER · EGRESS CONTROL

Data residency enforced; nothing leaves the region unbidden.



04 COMPLIANCE

PROOF & REPORTING

One definitive answer per agent action, ready for a regulator.



THE TEAM

BUILT BY INFRASTRUCTURE VETERANS.

Unicity Labs. PhDs in machine learning and cryptography with fifteen years building nation-state-grade security infrastructure for **DARPA, NATO, Lockheed, Verizon, and Maersk**. Now applying that to the rails your agents transact on.



Mike Gault

CEO

- PhD Electrical Engineering
- Built & exited Guardtime (ADX:IHC)
- Ex-MD, Barclays Capital



Tony Kenyon

CTO

- PhD Machine Learning
- 25 years shipping enterprise AI & infra
- Principal Architect: BT, Nokia, A10




Eric Leandri


CHAIRMAN

- AI Visionary
- CEO, Aleria (UAE Sovereign AI)
- 1,000+ customers · 250K agents live

WHERE WE OPERATE



CRYPTOGRAPHIC INFRASTRUCTURE TRUSTED BY



GET STARTED

SEE AN AGENT HIT A REGULATED BOUNDARY — **AND** **GET STOPPED.**

Watch an agent attempt a prescription, a cross-border transfer, an unconsented store – and get intercepted in-path, live, on your own infrastructure.

BOOK A LIVE DEMO →

info@unicity.ai · unicity.ai · +971 50 884 3756

APPENDIX

ADDITIONAL USE CASES.

DEEP-DIVES · SHADOW AI · DLP · COST CONTROL · ROUTING

USE CASE • VISIBILITY

AGENT SPRAWL & SHADOW AI.

You cannot secure what you cannot see. Shadow AI is present in ~40% of hospitals, and ~57% of professionals have used unauthorized tools – PHI pasted into a consumer LLM is an irreversible exposure legacy DLP and SIEM miss. Unicity is the mandatory registration and monitoring layer for all machine identity.

→ CRYPTOGRAPHIC IDENTITY


Every agent is assigned a revocable identity bound to a human principal. Eliminates static API keys.

→ GLOBAL INVENTORY

Map every active agent and its authorized scopes across the organization.

→ DRIFT CONTAINMENT

Instantly quarantine rogue agents that deviate from baseline execution patterns.



USE CASE • PHI EXFILTRATION PREVENTION

PHI EXFILTRATION PREVENTION.

Unicity sits directly in the transport flow, inspecting LLM payloads before they ever leave the hospital perimeter – so patient data can't be exfiltrated into an unauthorized model.

→ INLINE PHI REDACTION


Detect and mask PHI, PII, and identifiers in the prompt before inference ever runs.

→ GOVERNED PIPES

Sanctioned, monitored routes replace personal ChatGPT accounts and unmanaged tools.

→ COMPLIANCE LOGS

Cryptographic, tamper-evident logs of all patient data touched by agents, for HIPAA/GDPR audit.



USE CASE • COST CONTROL

COST OPTIMIZATION: EVERY AGENT GETS A BUDGET.

On thin hospital margins, a poorly engineered agent stuck in a hallucination cycle can burn thousands in token costs. We provide the kernel-level circuit breakers.

→ HARD RATE LIMITING


Enforce strict execution thresholds at the kernel level before cost is incurred.

→ TOKEN QUOTAS

Assign budget limits. The AOS Kernel blocks execution instantly when reached.

→ ROI ATTRIBUTION

Track exactly how much compute each autonomous workflow consumes.



COST OPTIMIZATION: INTELLIGENT ROUTING.

Every agent request is scored before it runs and directed to the right model – by task complexity, cost, latency, and security clearance. The routing decision is itself an inspected, recorded event, so traffic is matched to capability without ever escaping the security boundary.



— CAPABILITY-MATCHED

Heavy work reaches strong models; trivial work never pays for them.

— CLEARANCE-AWARE

PHI is only ever routed to a PHI-authorized target – never to a model that isn't.

— RECORDED

The routing choice is an inspected event – every decision is explainable after the fact.


SECTION

TECHNICAL APPENDIX.

AOS KERNEL · WHY AOS · SWAPPABLE CAPSULES · SEMANTIC INTERCEPT FABRIC

AOS: THE KERNEL

AOS is the secure enforcement kernel that tools like Claude Code and OpenClaw run **on top of**, not alongside. A user-space microkernel with syscalls, airlocks, and sandboxed isolation – the kernel stays minimal while everything else lives in user-space.



User-Space Microkernel

AOS models itself as an OS – with syscalls, airlocks, and user-space isolation. The kernel (aos9d) and system SDK (aos9-sys) are the only core components.

WASM Capsule Sandbox

All plugins run as WASM capsules inside a sandboxed runtime. OpenClaw plugins are compiled into capsules via aos9-openclaw, with host calls thunked into aos9::sys.

Enforcement Layer

Security interception, budget enforcement, and audit logging happen at the runtime level – below the agent, not beside it. Tools cannot bypass the kernel.

Everything Is User-Space

CLIs, OpenClaw, frontends, and even LLM providers are external user-space components communicating over IPC. The kernel stays minimal and ignorant of legacy ABIs.

WHY AOS

Existing approaches each solve part of the problem – system prompts control behaviour, containers isolate execution, policy engines enforce rules. But none deliver the full stack: **multi-user isolation, swappable components, audit, and sandboxing** in a single runtime. AOS gives you a pluggable brain in a padded room with a mail slot – the loop is **isolated from everything it controls**. It can't bypass the sandbox, skip the audit, or access tools it hasn't been granted.

APPROACH	SINGLE USER LOCAL DEV	MULTI-USER SHARED RUNTIME	SWAPPABLE LLM	SWAPPABLE AGENTIC LOOP	AUDIT TRAIL	EXTENSION	SANDBOXING
System prompt	Yes	No	No	No	No		No
Confirmation prompt	Yes	No	No	No	No		No
MCP + process isolation	Yes	Sort of	No	No	No		Partial
Policy engine (OPA/Cedar)	Yes	Yes	No	No	Yes		No
Container/VM	Yes	Yes	No	No	Partial		Yes
WASM plugins	Yes	Yes	No	No	No		Yes
AOS	Yes	Yes	Yes	Yes	Yes		Yes

APPENDIX - KERNEL

SWAP A CAPSULE. CHANGE EVERYTHING.

Because intelligence is no longer baked into the OS, anyone can change how AOS thinks without touching a single line of core code. **Swap a capsule, change the config, restart.** That's it.

RUN COMPLETELY OFFLINE

FINANCE, HEALTHCARE, DEFENCE

Swap the cloud provider capsule for a local one (Ollama, vLLM). Everything else keeps working. One-line config change, fully air-gapped.

NOVEL AGENT ARCHITECTURES

AI RESEARCHERS, STARTUPS

Plug in a debate system, MCTS planner, or chain-of-verification loop. The rest of the OS – tools, sandboxing, audit – works unchanged.

TRANSPARENT LLM CACHING

TEAMS RUNNING AGENTS AT SCALE

Drop in a caching capsule as middleware. 30-60% fewer API calls for repetitive workloads. No code changes, no redeployment.

AUTONOMOUS OVERNIGHT AGENTS

DEVELOPMENT TEAMS

Swap the orchestrator for an autonomous worker capsule. It writes code, runs tests, self-corrects – same sandbox, same audit trail.

MIX & MATCH PROVIDERS

TEAMS USING MULTIPLE MODELS

Install multiple provider capsules with a routing capsule. Direct each request to the best model by complexity, cost, or latency.

SHIP CUSTOM DISTROS

PLATFORM TEAMS, SAAS COMPANIES

One codebase, infinite configurations. Each customer gets a different capsule set. Security patches ship to everyone simultaneously.

SEMANTIC INTERCEPT FABRIC.

"the semantic firewall"

The enforcement point for **HIPAA, GDPR, and AI-Act** policy. Goes beyond **action-based** controls to include **content-based** ones – prompt injection, indirect injection, PHI leakage – with domain-specific clinical ontologies.

A LAYERED MODEL:

- 1. **Rule engine** – High-performance pattern matching, candidate identification, rule evaluation. Dynamic rule updates.
- 2. **ML Detection** – Low-latency inference and anomaly detection. Plugin-expandable (ONNX-compatible) model loading.
- 3. **DLP scanning** – For secrets / IP exfiltration – pattern matching plus name-entity recognition.

Rust-based implementation with optimised search and rule algos. Targets low-latency classification, depending on context (risk tolerance, industry sector), message type, and depth of content screening required.

